

PSTAT 135 Final Project

Aapthi Nagesh, Alex Roginski, Yuchen Wu, Nicholas Axl Andrian

1. Introduction

In modern American politics, advertisement to potential voters is essential for a political party's success at the polls. In our research, we attempt to utilize a variety of voter demographics, such as age and income level, to predict voter party affiliation. From this data, we create random forest and multinomial logistic regression models to predict party preference among voters. Then, we will analyze feature importance to see which predictors are the most influential on party membership.

1.1 Initial Setup

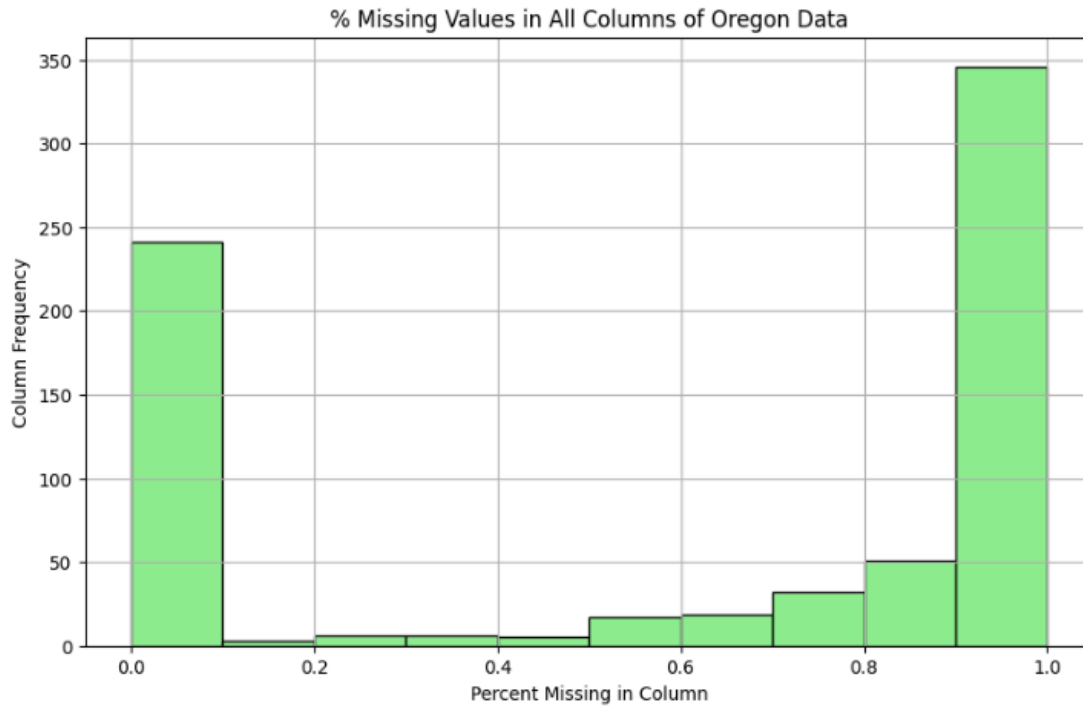
We will focus our investigation on the state of Oregon, whose voter file has over 2 million rows of data and 726 columns of predictors. Many of these predictors are irrelevant to our analysis and contain missing data as well, so we must clean our data set. After cleaning, this is what our dataset looks like:

county	county_ascii	county_full	county_fips	state_id	state_name	lat
Multnomah	Multnomah	Multnomah County	41051	OR	Oregon	45.5468
Washington	Washington	Washington County	41067	OR	Oregon	45.5601
Clackamas	Clackamas	Clackamas County	41005	OR	Oregon	45.188
Lane	Lane	Lane County	41039	OR	Oregon	43.9388
Marion	Marion	Marion County	41047	OR	Oregon	44.9033

only showing top 5 rows

1.2 Missing Data Analysis

Now that we have our voter data downloaded into a data frame, we can analyze the prevalence of missing data.



In this plot, each bar represents the number of columns with corresponding percentages of missing data. Out of 726 columns, 246 of them have 0-10% missing data, so we will choose our predictor variables from these. We will drop the remaining columns.

1.3 Column selection

We chose 15 predictor variables that we hypothesized would be important for voter party preference. Among these are voter age, gender, years of education, income, housing value, county, and several others. “Parties_Description”, or the party that each voter is affiliated with, is the outcome variable.

1.4 Merging voter data with country code dataframe

Now, we can join our two data frames, `df_counties` and `data`, to make the final dataframe that we can use throughout our analysis. We will join these on ‘counties’.

1.5 Binning party descriptions

The last step to cleaning our data is to put our party descriptions into bins. There are 10 unique parties listed in the voter files, but only a very small percentage of voters voted for a majority of the parties, so we combined those parties into a bin called ‘Other.’

```
+-----+
| Parties_Description|
+-----+
|      Republican|
|      Progressive|
|Registered Indepe...|
|              Other|
|      Libertarian|
|              Green|
|      Constitutional|
|Working Family Party|
|      Democratic|
|      Non-Partisan|
+-----+
```

After grouping the less common parties, this is what our outcome variable options look like:

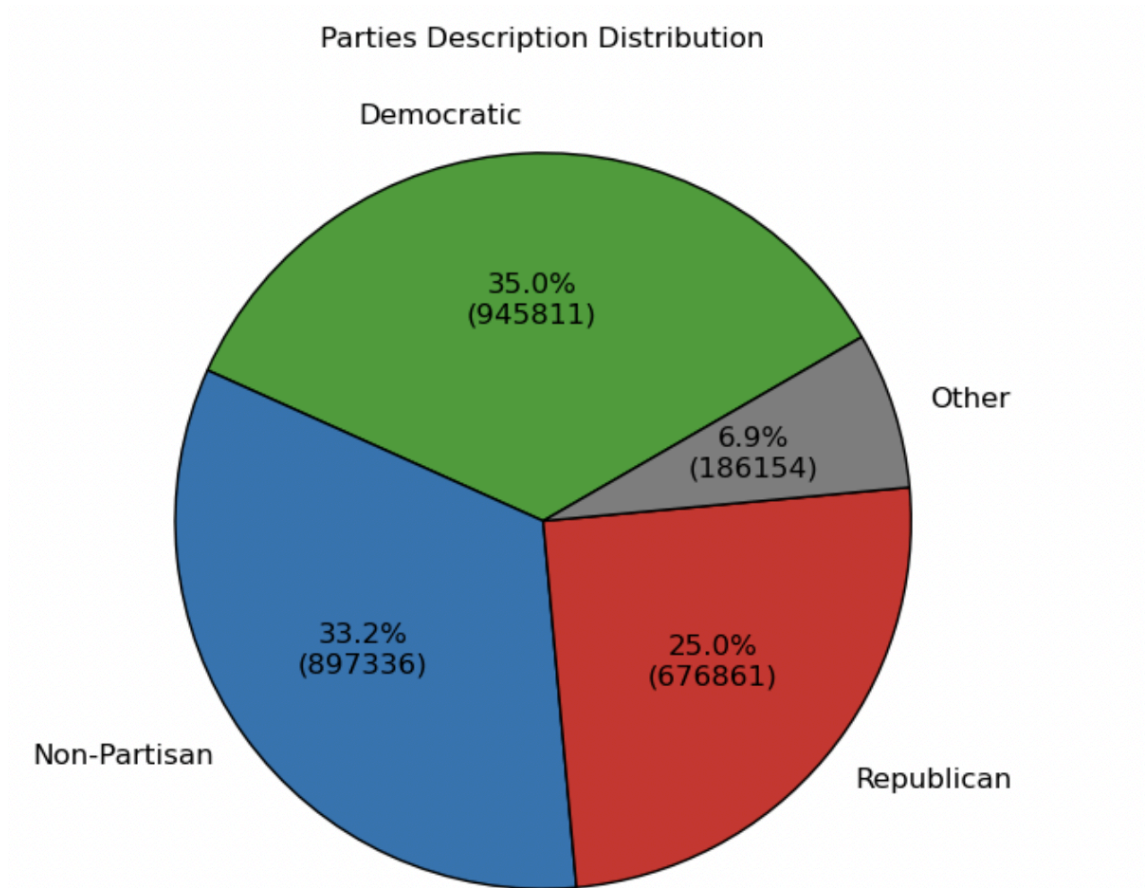
```
+-----+
|Parties_Description|
+-----+
|      Republican|
|              Other|
|      Democratic|
|      Non-Partisan|
+-----+
```

Now that we’ve prepared our data, we can begin our exploratory data analysis.

2. Exploratory Data Analysis

2.1 Party Preference Distribution

The first step of our exploratory data analysis is to look at the distribution of our response variable Party Description. To do this, we can look at a pie chart.

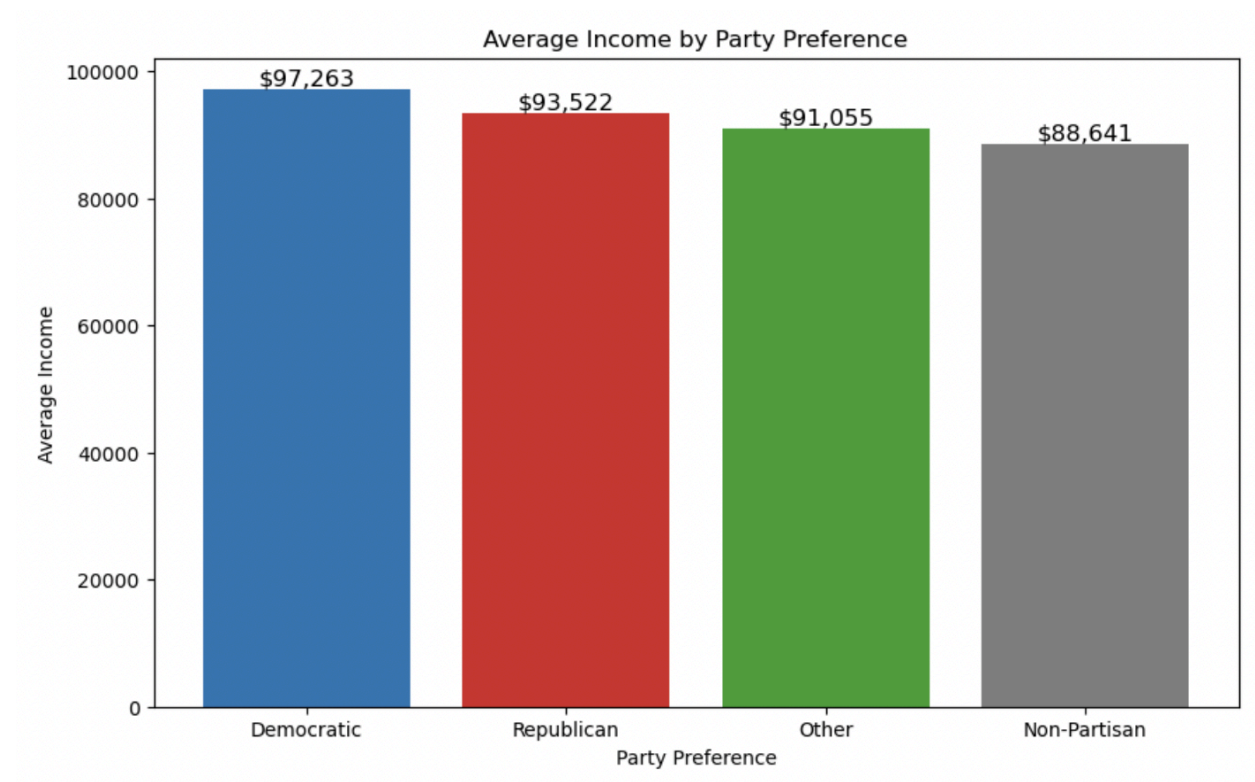


The Democratic party received the most votes, followed closely by the Non-Partisan and Republican parties. Since the rest of the parties made up less than 7% of the data, they've all been binned into the 'Other' category. It's important to note that this distribution is specific to the Oregon voter files, not the whole country. We found it surprising that the Non-Partisan voter population is significantly larger than the Republicans.


2.2 Income vs. Party Preference

In this section we will examine the relationship between economic status and political party preferences in Oregon. As we know, the state of Oregon tends to lean Democratic in its elections, so it is a particularly interesting case study for examining and analyzing the dynamics between economic status and party preferences. Specifically, we will be examining whether household

income levels influence people's party preferences in Oregon, and how party preferences vary across different income levels.



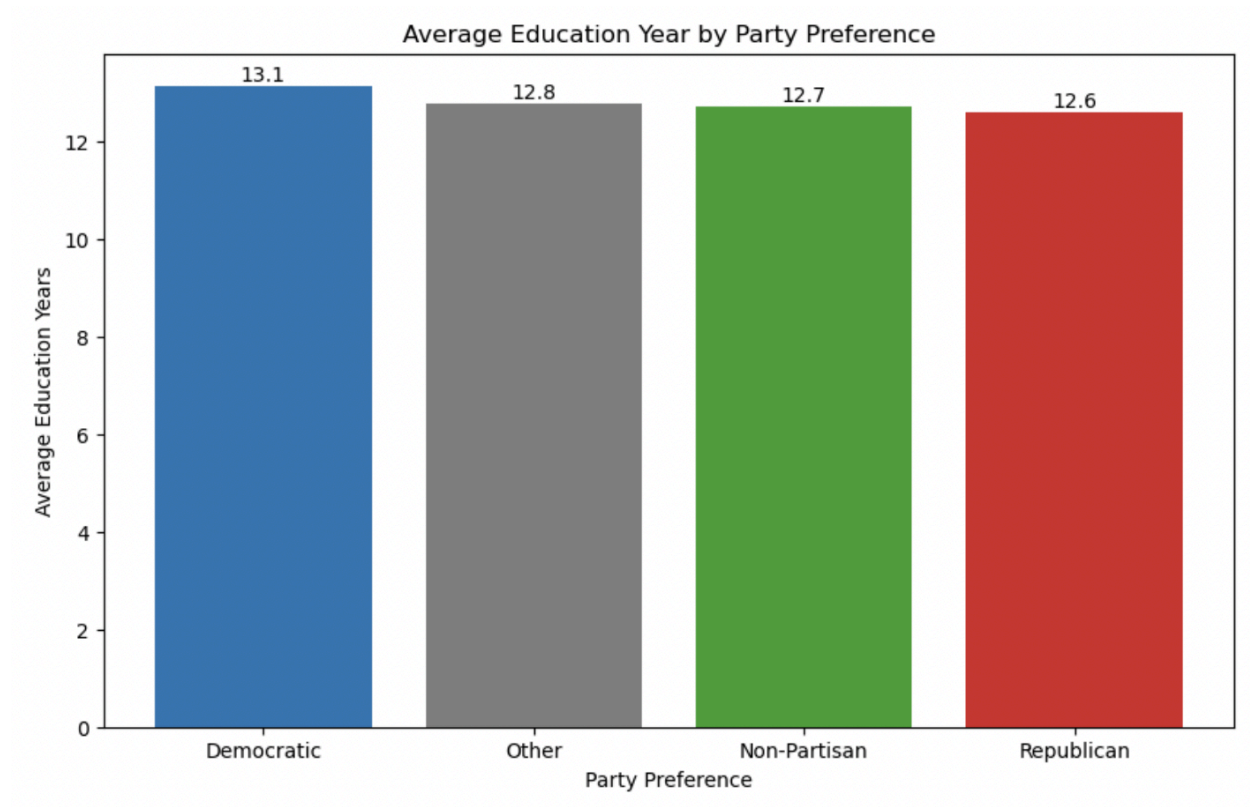
The bar plot above shows the average household income for each party preference. We can see that Democrats have a slightly higher household income on average than members of other parties. Non-Partisan voters have the lowest household income on average. Overall, the differences between parties are minimal, we should take a further look by breaking income into distinct income levels, and see which party each level tends to favor. Here we use [data](#) provided by the state of Oregon to classify voter incomes into levels accordingly: Less than 40k a year is Low, 40k - 100k is Medium, and over 100k is High.



From the plots, we can see that the Democratic Party is the preferred party in both low-income and high-income households, with them leading the high-income households by a significant margin (~7%). Preferences for the Republican party are fairly stable across all income levels, potentially indicating a consistent base of support regardless of income. A significant portion of each income level also prefers the non-partisan party, with the highest in the medium-income house in which they lead the group by a small margin (~3%). These results suggest that income level does correlate with political party preference in the state of Oregon to an extent, however, the relationship seems very weak.

2.3 Education Level vs Party Preference


The next section focuses on the highest level of education of the voters in Oregon. We aim to see whether the highest average education level varies among voters of different party preferences. This can help us determine if education levels are significantly influential in predicting party affiliation amongst voters in Oregon.



Here, we can see that average years of schooling do not have a large impact on voting preferences. The average education level of Oregon voters is between 12 and 13 years. Democratic voters have slightly more years of schooling than the other parties, but only by a very small margin. Republican voters have slightly less schooling than the other parties, but also by a very small margin. Similar to the previous section, we should take a closer look by classifying education years into three distinct groups: Low, Middle, and High. Since not much is known about education years, let's take a look at the max and min values for it to determine the range we should set each group to.

```
+-----+-----+
|Max Education Years|Min Education Years|
+-----+-----+
|           17|           10|
+-----+-----+
```


Max and Min education years in Oregon data are 17 and 10 years respectively. With this knowledge, we will assign anyone who has below 12 years of education to the Low group, 12-15 years to the Middle group, and 15+ years to the High group. Below is a visualization of the groups and their party preferences.



From the plots, we can see that the Democratic Party is the clear favorite amongst voters who have a Middle or High education level, and they lead the High education level by an enormous margin over the other parties. However, they have relatively low support from voters who have a Low education level. Non-partisan voters, on the other hand, have the most support from voters with Low education levels compared to the other parties, and the support decreases relatively as we move up the educational ladder. Support for the Republican party is also mainly concentrated in the Low education level, and their support is very weak amongst the Highly educated. These pie charts do suggest that Education level plays a significant role in people's party preferences, as seen by the large differences.

2.4 Property Type vs Party Preference

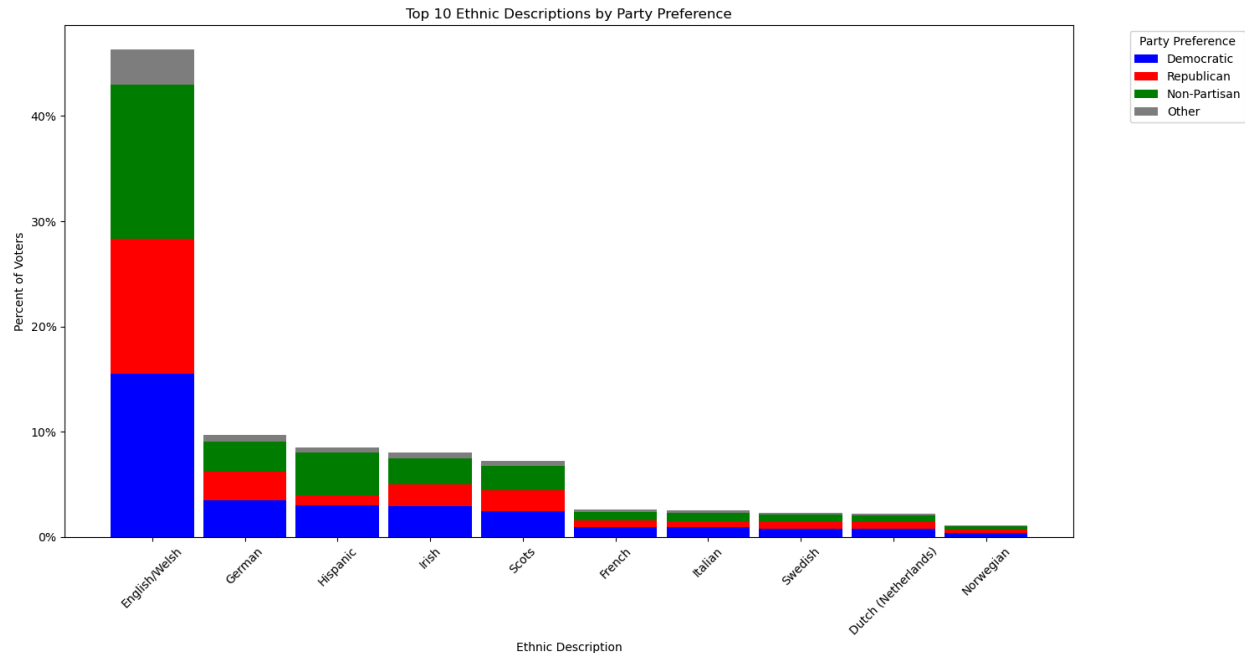
A previous section has suggested there may be a relationship between household income and party preferences, so in this section, we start to investigate if different property types owned – residential, apartment, condominium, etc. – are related to political party preferences. The variable 'CommercialData_PropertyType' in our dataset categorizes the type of property owned by each individual voter.



From this stacked bar plot, we can see that for all Oregon voters, regardless of party preference, the residential property type is the highest reported. Republicans have the most support from voters who live in a residential property, while voters who live in an apartment tend to favor Non-Partisan and Other parties more. However, the differences are not very large, so we cannot conclude property type is an influential indicator of party preference.

2.5 Top 10 Ethnic Descriptions vs Party Preference


While exploring the ethnic descriptions predictor, we decided to narrow our EDA to the top 10 ethnicities that are the most prevalent in Oregon. This decision was made when we saw that the large majority of our voters fall into these 10 ethnicities. By narrowing down the ethnicities, we can still get a thorough understanding of the distribution and its interaction with party affiliation.



While looking at the distribution of ethnicity in our data, we can see from the bar plot that by far, the largest amount of voters are English/Welsh. This is worthy of note since Oregon is known to be less diverse in its population than other, more populated states. The parties are split relatively evenly among the top 10 ethnic descriptions, with the exception of the Other bin, which has the least votes in all ethnic groups. Also, we can see that the Hispanic ethnic population of Oregon has a smaller amount of Republican voters and more Non-partisan voters in its distribution than the other ethnicities.

2.5 Age vs Party Preference

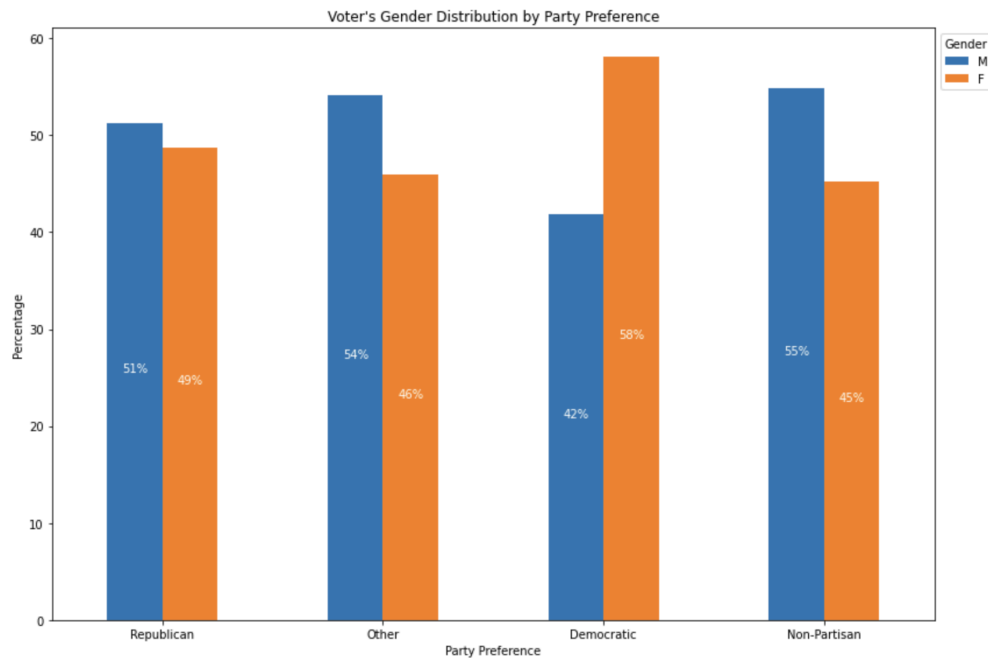
Next, we can look at the age distribution of the voters in our data and the interaction between it and party affiliation. To understand the distribution of ages better, we binned our data into three levels and used this to look at different age groups separately and study their preferences. We binned the ages into an 18-35 age group, a 36-64 age group, and a 65+ age group.



For the distribution of party preferences regarding different age groups, we can use pie charts. In the first and youngest age group, Non-Partisan votes make up nearly half of the distribution, followed by Democratic votes, which make up about a third. In the second age group, Democratic, Non-Partisan, and Republican votes are more evenly split, with Democratic votes taking the majority this time. Finally, for the third age group, we can see that the Non-Partisan votes are much smaller in percentage and there are both more Democratic and Republican votes in the distribution. This indicates that younger voters in Oregon tend to vote more in the middle, but they also tend to choose a side ideologically as they get older, and the vote splits more between the Democratic and Republican parties.

2.6 Gender Description vs Party Preference

The final section of our EDA is focused on the gender description distribution of Oregon voters and its interaction with party preferences.



Here we can see that for all parties, males in Oregon have a higher percentage of the votes, with the exception of the Democratic party where females make up 58% of the votes while males only make up 42%. This shows that gender does seem to be somewhat significant when it comes to party preference in Oregon, so we can study this further in our modeling.

3. Modeling

Now, we will create a random forest and multinomial logistic regression to predict voter party preference.

3.1 Preparations for modeling

To prepare the data for the models, we must change the data formats of some of the columns in the voter data. Some of the numerical data, such as home value, is stored in the data with dollar signs, e.g. \$150,000, so we have to convert it to numeric values. Also, all of the numeric columns, such as voter age, are also stored as categorical columns, so we have to convert them to numeric.

Data_EstimatedHHIncomeAmount	CommercialData_EstimatedAreaMedianHHIncome	CommercialData_AreaMedianHousingValue	CommercialData_AreaMedianEducationYears
60562.0	60562.0	255416.0	12.0
60562.0	60562.0	255416.0	12.0
60562.0	60562.0	255416.0	12.0
60562.0	60562.0	255416.0	12.0
60562.0	60562.0	255416.0	12.0
60562.0	60562.0	255416.0	12.0
138000.0	60562.0	255416.0	12.0
138000.0	60562.0	255416.0	12.0
60562.0	60562.0	255416.0	12.0
60562.0	60562.0	255416.0	12.0

After the data preparation, this is how our columns look:

```
[('County', 'string'),
 ('Mailing_Addresses_City', 'string'),
 ('Ethnic_Description', 'string'),
 ('Voters_Gender', 'string'),
 ('Voters_Age', 'float'),
 ('Mailing_Families_HHCount', 'float'),
 ('Mailing_HHGender_Description', 'string'),
 ('Residence_Families_HHCount', 'float'),
 ('CommercialData_EstHomeValue', 'float'),
 ('CommercialData_EstimatedHHIncomeAmount', 'float'),
 ('CommercialData_EstimatedAreaMedianHHIncome', 'float'),
 ('CommercialData_AreaMedianHousingValue', 'float'),
 ('CommercialData_AreaMedianEducationYears', 'float'),
 ('CommercialData_PropertyType', 'string'),
 ('CommercialData_StateIncomeDecile', 'string'),
 ('Parties_Description', 'string'),
 ('county_fips', 'int')]
```


We see that the columns have been coerced into formats necessary for modeling.

3.2 Random Forest Model

After fitting the model on the training data and evaluating the model on the testing data, we got an accuracy score of 0.459. Generally, above 70% is considered a valuable and realistic model. Our low accuracy score suggests that our model did not effectively capture the factors that predict a voter's party preference.

3.3 Feature Importance (Random Forest)

Using our Random Forest model, we next looked at what features were most important to the model.



The Random Forest feature importance analysis shows that voter age, the median number of years of education in a voter's area, and their county are the most important variables among those we studied in predicting a person's voting preference. Less important is the type of property a voter lives in, the voter's gender description, and their ethnic description. Finally, the least important features are the estimated area median income, the total number of voters residing at the given residence address, estimated home value, the state income decile, and estimated income for each voter.

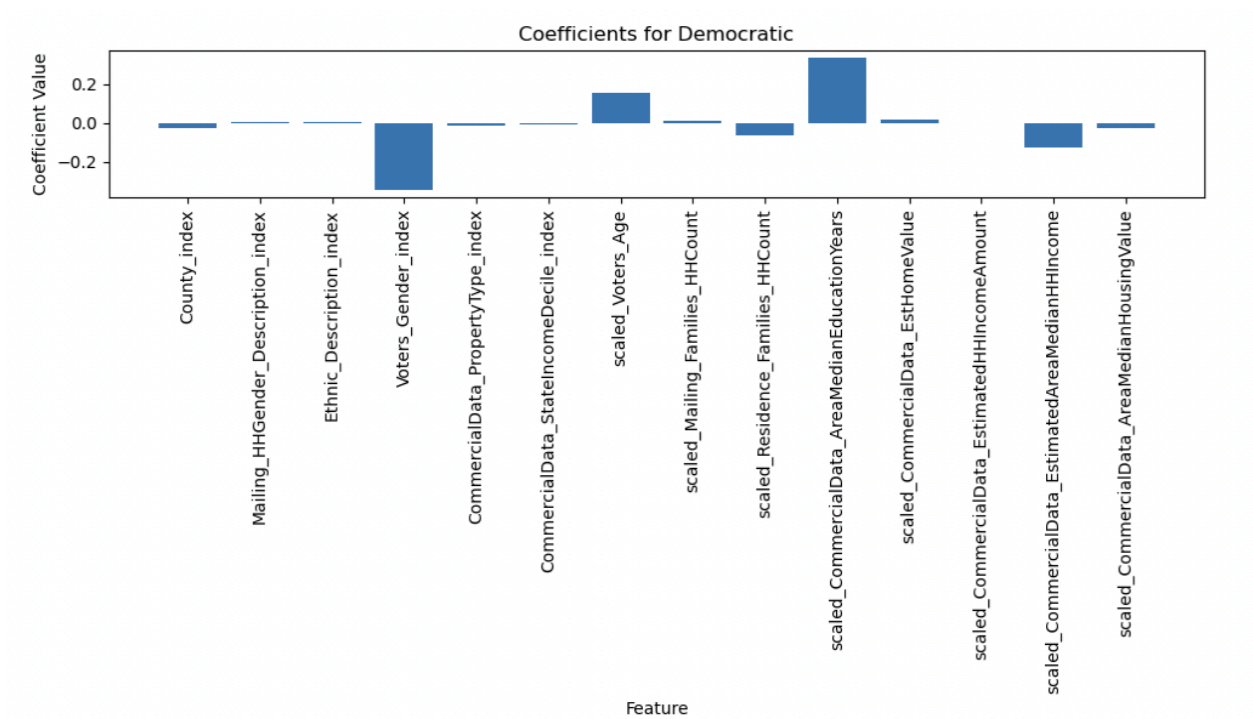
3.4 Multinomial Logistic Regression

Next, we're going to fit a multinomial logistic regression to see how it compares with a Random Forest.


After fitting the model on the training data and evaluating the model on the testing data, we got an accuracy score of 0.461. It performs about the same as the random forest model with a difference in accuracy of 0.002 between the models. This suggests that both models perform very poorly. The low accuracy score indicates that neither model could effectively use the input variables to predict voter party preference.

Now, we will look at the feature importance for the multinomial logistic regression to see which variables are most important for which parties.

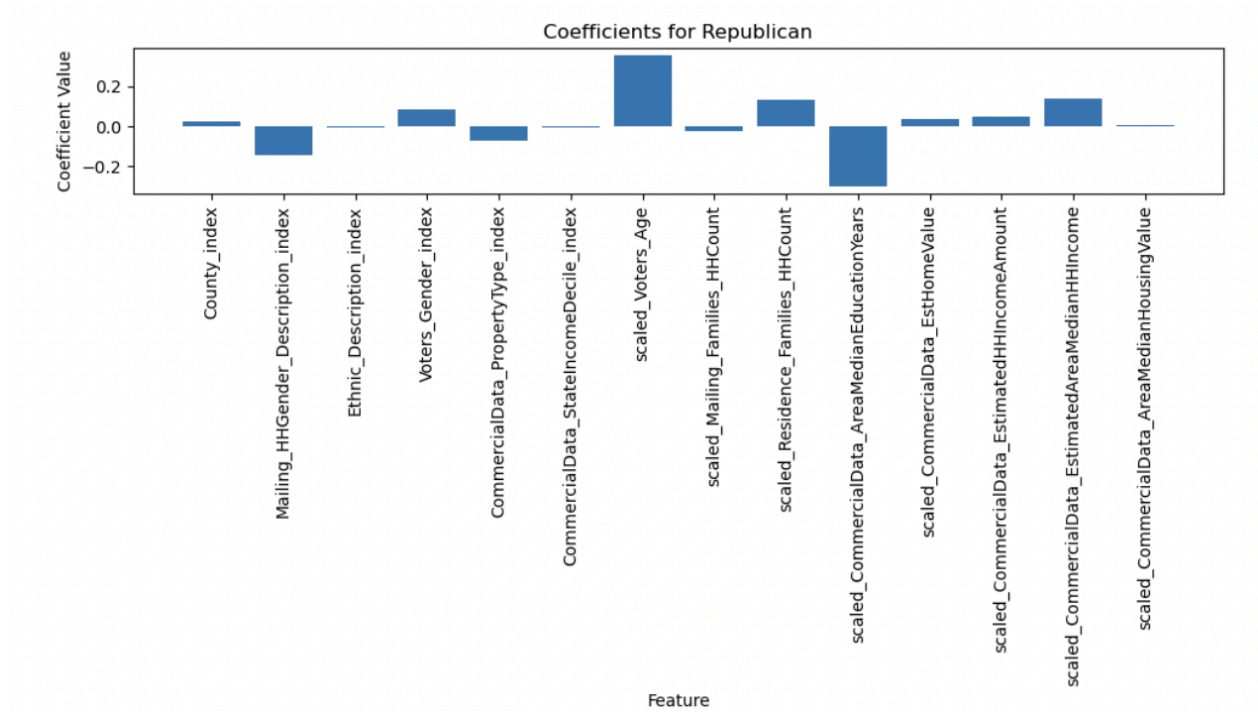
3.5 Feature Importance (Multinomial LR)




Looking at the Democratic feature importance plot, we see that being female is associated with being a member of the Democratic party. We also see that older voters tend to associate more with the Democratic party. Finally, having a higher number of education years is associated with being part of the Democratic party. The other variables do not seem to have much of a strong correlation with being part of the Democratic party.



The Non-Partisan feature importance plot suggests that being younger is the strongest indicator of being non-partisan. It also skews voters towards being male. The other variables are less important in this model.



The Republican feature importance plot suggests that being older and having fewer years of education is associated with being part of the Republican party. Also, having a higher area median income is associated with the Republican party. The other variables are less important. Gender isn't as important for predicting Republican party membership as for Democrats.



Being male is associated with being part of the other parties besides Democrat and Republican. Also, being younger is associated with being part of other parties as well. The rest of the variables are not as important in this feature importance plot.

Here is the top input for each of the top 5 predictors given a party preference:

```

Top 5 predictors for Democratic:
scaled_CommercialData_AreaMedianEducationYears    0.335227
scaled_Voters_Age                                  0.153662
scaled_CommercialData_EstHomeValue                 0.023586
scaled_Mailing_Families_HHCount                   0.014983
Mailing_HHGender_Description_index                 0.005224
Name: Democratic, dtype: float64

Top 5 predictors for Non-Partisan:
Voters_Gender_index                                0.119641
Mailing_HHGender_Description_index                 0.112415
CommercialData_PropertyType_index                 0.054593
scaled_Mailing_Families_HHCount                   0.035982
scaled_CommercialData_EstimatedAreaMedianHHIncome 0.030349
Name: Non-Partisan, dtype: float64

Top 5 predictors for Republican:
scaled_Voters_Age                                  0.355064
scaled_CommercialData_EstimatedAreaMedianHHIncome 0.135777
scaled_Residence_Families_HHCount                 0.134537
Voters_Gender_index                                0.086830
scaled_CommercialData_EstimatedHHIncomeAmount     0.048444
Name: Republican, dtype: float64

Top 5 predictors for Other:
Voters_Gender_index                                0.138483
scaled_CommercialData_AreaMedianHousingValue      0.036660
CommercialData_PropertyType_index                 0.030860
Mailing_HHGender_Description_index                 0.027691
scaled_CommercialData_AreaMedianEducationYears    0.009987

```

Here is the top input for each predictor within each party group (standardized):

```
Top inputs for Democratic:
Voters_Gender_index: 0.0
scaled_CommercialData_AreaMedianEducationYears: [-0.7529937387323782]
scaled_Voters_Age: [1.029567987522468]
scaled_CommercialData_EstimatedAreaMedianHHIncome: [0.16125966517928514]
scaled_Residence_Families_HHCount: [-0.04601681339265395]

Top inputs for Non-Partisan:
scaled_Voters_Age: [-1.4801129341029797]
Voters_Gender_index: 1.0
Mailing_HHGender_Description_index: 0.0
CommercialData_PropertyType_index: 0.0
scaled_Residence_Families_HHCount: [-0.04601681339265395]

Top inputs for Republican:
scaled_Voters_Age: [0.7091831890170917]
scaled_CommercialData_AreaMedianEducationYears: [-0.7529937387323782]
Mailing_HHGender_Description_index: 0.0
scaled_CommercialData_EstimatedAreaMedianHHIncome: [0.16125966517928514]
scaled_Residence_Families_HHCount: [-0.04601681339265395]

Top inputs for Other:
Voters_Gender_index: 1.0
scaled_Voters_Age: [-0.9995357363449153]
scaled_CommercialData_EstimatedAreaMedianHHIncome: [0.16125966517928514]
scaled_CommercialData_AreaMedianHousingValue: [1.3582325608122368]
CommercialData_PropertyType_index: 0.0
```

Finally, here is the top input for each predictor within each party group (True values). It's important to note that the true value conversions were hard coded for this run.

```
Top inputs for Democratic:
Voters_Gender: F
CommercialData_AreaMedianEducationYears: 12.18
Voters_Age: 52.88
CommercialData_EstimatedAreaMedianHHIncome: 59024.02
Residence_Families_HHCount: 2.1

Top inputs for Non-Partisan:
Voters_Age: 43.4
Voters_Gender: F
Mailing_HHGender_Description: Mixed Gender Household
CommercialData_PropertyType: Residential
Residence_Families_HHCount: 2.06

Top inputs for Republican:
Voters_Age: 50.86
CommercialData_AreaMedianEducationYears: 12.43
Mailing_HHGender_Description: Mixed Gender Household
CommercialData_EstimatedAreaMedianHHIncome: 76699.16
Residence_Families_HHCount: 2.23

Top inputs for Other:
Voters_Gender: M
Voters_Age: 37.05
CommercialData_EstimatedAreaMedianHHIncome: 74907.36
CommercialData_AreaMedianHousingValue: 191123.07
CommercialData_PropertyType: Residential
```

Here we see the most frequent categorical inputs and the mean numerical inputs for the top 5 features with respect to each party preference. The data comes from the predictions of our logistic model, suggesting that these characteristics are most common for members of each party preference. Political parties can use this information to advertise to potential voters to gain more votes. For example, if advertisers from the Democratic party are trying to get more Democratic votes, our model says they should advertise to women around age 53 that have a median education of 12. However, note that due to low model accuracy, the model's suggestions are prone to being mistaken.


We found that running the multinomial logistic regression multiple times yielded vastly different results for which attributes were associated with which parties. This might happen because our accuracy is low, so there is more variance in the model each time we run it.

4. Generalizability

Now that we've fit the random forest and multinomial logistic models, let's see how well they do on other datasets. We chose to test the two models we trained on Oregon data on three other

states. We chose Washington since it is Oregon's neighbor, so we infer similar accuracy since voter behavior might be similar in neighboring states. Our second state was Maine because it's far away from Oregon and we were wondering if the model would work in different states. Our last selected state is Louisiana since it is far away from Oregon and also very culturally different since it's in the Deep South.

We will use each state's data as testing data for the model trained on Oregon and compare accuracies between states and between models.



In terms of comparison between random forest and multinomial logistic regression, we can see that they perform very similarly regardless of the state's data they're tested on. Also, we see that while the random forest and multinomial logistic regression perform about the same for the Oregon data, the logistic model outperforms the random forest by about 1% accuracy for the

other states. This might suggest that the logistic model does better applied to new data (although a 1% accuracy difference might not be considered significant).

While comparing between states, we see that accuracy for both models is highest in Washington. This makes sense since Washington is physically closest to Oregon. Maine, which is across the country from Oregon, scores about 4% lower in accuracy. This might be because Maine is so far from Oregon. Louisiana has the lowest accuracy (6% lower than Washington), which makes sense since Louisiana is physically distant and the population attributes are quite different (income, party preference distribution, etc.).

In general, we can conclude from these findings that our model trained on the Oregon dataset, has an accuracy that decreases depending on its cultural and physical distance from Oregon. This makes sense as different voter attributes might be more important in other states, be it income, family size, etc. This can cause them to have more significant predictor coefficients that are not necessarily captured as well by the ones we have seen in Oregon's feature importance results.

5. Conclusion

In this project, we looked at what predicts voter party preference in the state of Oregon. For our Exploratory Data Analysis, we chose variables expected to impact voter party preference, and we explored and visualized the relationships. What we found from the graphs is that economic status, measured by household income levels and property type owned, generally displayed a weak relationship with party preferences, suggesting little to no impact. This was a little surprising as economics and taxes are usually a hot issue in political debates. On the contrary, educational level and age, each classified into their own distinct groups based on a criteria, displayed a clear difference in preferences, suggesting age and educational level plays a significant role in people's party preferences.

Next, we fitted models to the data. By fitting a random forest and multinomial logistic regression to 15 predictor variables, we attempted to evaluate what factors are most important in a person's party preference. Overall, both models did poorly with random forest and logistic accuracies of 0.459 and 0.461 respectively. This suggests that the models cannot accurately predict voter party preference in Oregon given the chosen variables.

Our random forest model found that a voter's age, their median number of years of education in their area, and their county are the most important variables among those we studied for predicting voter party affiliation. Less important variables included area median income, estimated home value, and state income decile.

Our multinomial logistic regression analyzed which factors are most important for each party. Some of its findings included that being female is associated with being part of the Democratic party, younger people tend to be part of “Other” parties, and that being older is associated with being Republican.

From our logistic regression predictions, we calculated the attributes of the average member of each party. This data can give campaigners for the Democratic party insight on which kinds of people to target for advertising. However, due to low model accuracy, the averages are often incorrect or misleading. We also found that while running our models, we would get vastly different results for which attributes were associated with which parties. This might occur due to low accuracy of the models.

We also found that our model is not very generalizable. After evaluating voting records from Washington, Maine, and Louisiana as testing data for our random forest and logistic regression trained on Oregon data, we found poor accuracy scores of around 0.42, 0.38, and 0.37, respectively for both models (with the logistic model having a 1% accuracy increase compared to the random forest). This suggests that the model does worse with increasing physical distance and difference in voter behavior and lifestyle. Also, note that even if the states we tested on remained very close in accuracy scores to the original Oregon models, the accuracy is still far too low to make meaningful conclusions.

In future research, we would like to explore how we can use the model data to convert potential voters into a party preference. Using the top input per most important feature column results, we calculated which factors are most important to each party, but low accuracy from the model gives us low confidence in these findings. Also, the models trained on only Oregon data struggle to accurately predict voter preference in other states, as well.

One important question we’d love to answer in further research is how would we tune our parameters better to have a better fitting model? An approach that we did attempt was the hypertuning of the parameters which unfortunately resulted in an unreasonably long loading time before returning a timeout error. We could however re-evaluate the columns we chose to work with and the criterias we ended up choosing from. Another approach we could look at would be to instead survey a group of a certain party preference in real life to conduct an investigation on whether or not it fits the predictions we have gotten from our models. All these potential improvements allow our project to have a wide room for improvement.